THE AI MIRROR HOUSE

Same reflection. Different filters. Better ideas.

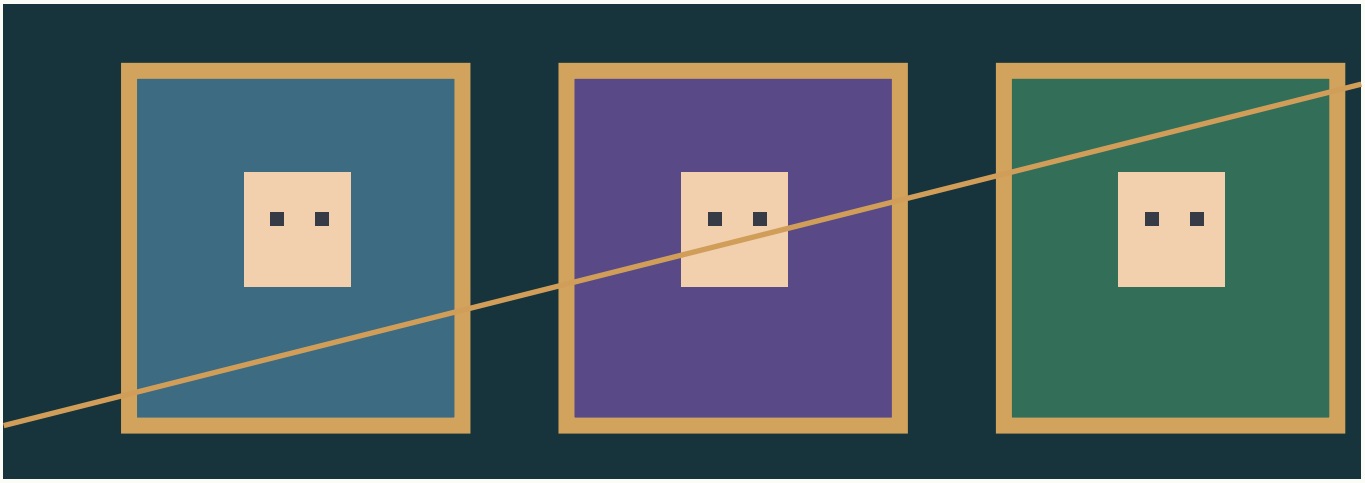

Aletheia Protocol (https://github.com/KarstenEvans/aletheia-protocol) provides provenance, evidence, uncertainty, correction and human control. Thalia Protocol (https://github.com/KarstenEvans/thalia-protocol/blob/main/THALIA_PROTOCOL.md) adds measured humour and a little warmth, without ever replacing the facts.

An Aletheia editorial inspired by Jonathan Parsons' LinkedIn discussion (https://lnkd.in/p/dvU-Nkkw) about specialist AI tools.

The familiar face in the familiar mirror

Every morning, you look into the same mirror. The lighting may change. Your hair might be more adventurous than you intended. But the face looking back is recognisably yours.

AI can feel rather similar. Ask it for a LinkedIn article and it produces a polished article. Ask it to make the article more engaging, and before long everything becomes groundbreaking, every change is transformative, and someone is once again unlocking the future.

The prose is tidy. It may even be useful. But it starts to look as though the same invisible copywriter has been hired by half the internet.

A recent LinkedIn post by Jonathan Parsons makes a sensible argument: stop treating one general-purpose AI as the answer to every task. Use specialist tools when their purpose-built workflows do the job better. A presentation editor is not the same thing as a research notebook; a video editor is not a writing partner.

I agree with the problem. I am less convinced that changing applications, on its own, is the solution.

A different mirror can still show you the same face.

Add some colour. Then question the reflection.

Suppose you move from one mirror to another. One is Victorian and gilded. Another is a three-panel dressing-table mirror. The next is an enormous disco ball, offering 200 tiny opinions about your haircut.

Add makeup, coloured lighting, an outrageous hat or a new pair of glasses and you will see something different. But the real change comes when you start moving around, looking from another angle, and noticing something you previously missed.

That’s a better way to think about AI applications. Their interfaces and specialist capabilities are useful filters, not automatic guarantees of originality. Switching models may help, but so can switching the question, adding an unexpected example, changing the intended reader, or asking for the strongest case against your own argument.

This is more than a matter of taste. In a 2025 study of 2,200 college admissions essays, researchers found that GPT-4-generated writing contributed less diversity of ideas across a group than human-written essays, even where individual AI-assisted text could be strong. A 2026 review covering 19 studies found a small but statistically significant tendency towards homogenisation in human-AI creative work. The degree varied with the task. Neither study proves that every AI output is generic, or that changing tools cannot help. They do suggest that how we work with AI matters. [1] [2]

Mirror Filters: five ways to change the conversation

Rather than asking, ’Make this better’, try putting a few intentional filters in front of the mirror:

The Reader. Who will read this, and what might they misunderstand, doubt or genuinely care about?

The Other Angle. Give me a genuinely different interpretation, not the same paragraph wearing a new tie.

The Evidence. Which claims are supported? Which need checking? Which are opinions dressed up as facts?

My Voice. Keep the odd little details, the original experience and the phrases that actually sound like me.

The Pressure Test. What is the strongest reasonable objection to this idea? What would change my mind?

Aletheia calls this Reflect, Refract, Reality-check: first understand the original thought, then explore materially different directions, then test what remains standing. The aim is not to produce five decorative rewrites. It is to find something worth saying.

A trio in the dressing-table mirror

There is another rather delightful possibility. Think of a three-panel mirror where one singer tilts the side panels and suddenly becomes a trio. The voices appear to answer one another. Now imagine doing that deliberately with AI.

Start with one model and ask for the best explanation of your idea. Give its answer to a second model and ask it to challenge the assumptions and find missing perspectives. Take the critique to a third model and ask for an original alternative, grounded in the same facts. Return to the first, give it the disagreements, and ask for a more considered synthesis.

The mirrors are now talking to one another. AI gossip, if you like, but preferably with footnotes.

This is the spirit of Aletheia Go Walkabout: explore adjacent fields and unusual connections, return with possibilities, and distinguish intriguing speculation from verified evidence. The point isn't to let three chatbots congratulate one another until the biscuits run out. It's to make disagreement productive. [3]

There are important limits. Different models may repeat the same errors, cite one another's inventions, or blur where a claim originated. Keep the original brief and source links with every handoff. Do not send personal or confidential material between services without considering the privacy implications. Ask each AI to challenge, not merely approve. And give the human the final decision.

Where Aletheia and Thalia belong

This is where the two ideas can work together.

Aletheia is the anchor. It keeps the project context, evidence trail, alternatives, corrections and human judgement portable, rather than trapped in whichever chat window happened to generate the first draft.

Thalia is the creative companion. It can introduce warmth, play, humour and unexpected associations, provided it doesn't turn a factual question into a stand-up routine. A joke can open a door to a new idea. It cannot authenticate a source. The clown nose is not a peer-review certificate.

That suggests a practical experiment for Aletheia Writing Vibe: add an optional Mirror Filters mode for the user's own draft, while keeping its existing evidence-based author-analysis function separate. It could prepare a portable set of prompts and comparisons for whichever AI the reader prefers. Aletheia Learn could teach the same technique as a short, free-first exercise. These are proposed enhancements, not features already operating automatically.

A mirror, not a master

Perhaps the mistake is believing that the quality of an AI answer lies entirely inside the AI. The human supplies curiosity, purpose, knowledge of the world and the right to say: 'No, that sounds nothing like me.'

A specialist tool is sometimes exactly what a task needs. At other times, the best improvement is to ask a better question, invite a dissenting reflection, or walk away from the mirror long enough to have an original thought.

You don’t need eighteen mirrors to become a more interesting person. You need to know when to change the light, when to turn the mirror, and when to trust your own eyes.

So, when you next ask AI to improve your writing, don’t stop at ’Make it more engaging’. Try: ’Go Walkabout. Find a different angle. Challenge my assumptions. Keep my voice. Show me what you can prove.’

And if three AI models begin agreeing rather enthusiastically, ask whether they’ve actually checked the facts, or simply formed a very polite choir.

Sources and useful links

[1] Moon, K., Green, A. E., and Kushlev, K. (2025). Homogenizing effect of large language models (LLMs) on creative diversity. Computers in Human Behavior: Artificial Humans. https://doi.org/10.1016/j.chbah.2025.100207

[2] de Rooij, A., and Biskjaer, M. M. (2026). Does generative AI make us think alike? A systematic review and meta-analysis of homogenisation effects in human-AI co-creation. Behaviour & Information Technology. https://doi.org/10.1080/0144929X.2026.2726451

[3] Aletheia Learn, Go Walkabout: https://github.com/KarstenEvans/aletheia-app/blob/main/aletheia-learn/aletheia-learn-ai-fluency.md

Original LinkedIn discussion: https://lnkd.in/p/dvU-Nkkw

Aletheia Writing Vibe: https://karstenevans.github.io/aletheia-app/aletheia-writing-vibe/aletheia-writing-vibe.htm

Aletheia Protocol: https://github.com/KarstenEvans/aletheia-protocol

Thalia Protocol: https://github.com/KarstenEvans/thalia-protocol

Aletheia Go: https://karstenevans.github.io/aletheia/aletheia-go/

Aletheia Worldwide AI: https://karstenevans.github.io/aletheia/

Swindon.org.uk: https://swindon.org.uk/

Editorial and illustrations Aletheia AI, 20 October 2026. Original analysis, not an endorsement by the authors or companies discussed. Application features may change. Research claims were checked against the cited publications on 20 October 2026. Human review before publication is recommended.

MORE FROM ALETHEIA

Aletheia asks what is true. Thalia asks what is amusing and how to create humane original humour. GoWalkabout takes a bounded creative detour, then checks the ideas before returning.

GO WALKABOUT: https://karstenevans.github.io/aletheia/aletheia-go/go-walkabout.htm

ALETHEIA: https://github.com/KarstenEvans/aletheia-protocol

THALIA: https://github.com/KarstenEvans/thalia-protocol